

DB2 Evidence Loader

Python tool to load YAML evidence files into a DB2 database.

NOTES:

- With minor changes, it should also be applicable for other database providers.
- In the guide we use Db2 for Linux, UNIX, with minor changes it should also be applicable for Db2 for z/OS.
- The Python source code, DB2 schema, and configuration files are mentioned solely as example use cases. They should not be interpreted as any form of design commitment.
- This this sample code that may contains issues. Not to use in production.

Table of Contents

- Architecture
- File Structure
- Database Schema
- Database Setup
 - Step 1: Create Database (if needed)
 - Step 2: Create Schema
 - Step 3: Create Tables
 - Step 4: Verify Tables
- Installation and Configuration
 - Option 1: IBM DB Driver (ibm_db)
 - Option 2: JDBC Driver (JayDeBeApi)
- Usage
 - Configuration File
 - Using Environment Variables
 - Using IBM DB Driver
 - Using JDBC Driver
- Querying Samples
- Driver Comparison
- Additional Resources
- License

Architecture

The project uses a **Template Method** pattern with:

- `DB2EvidenceLoaderBase`: Abstract class containing all business logic
- `DB2EvidenceLoaderIBM`: Implementation using `ibm_db` driver
- `DB2EvidenceLoaderJDBC`: Implementation using `JayDeBeApi` (JDBC) driver

Artifact Uniqueness

Important: Artifacts are identified by their PATH (unique constraint):

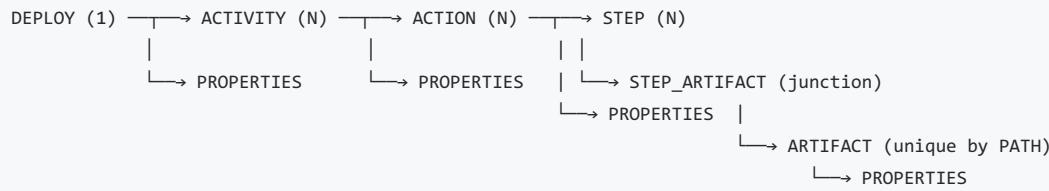
- Same artifact path = same artifact (reused across multiple steps)
- Many-to-many relationship between STEP and ARTIFACT via `STEP_ARTIFACT` junction table
- When an artifact with the same path is encountered, the existing artifact is linked instead of creating a duplicate
- The `ARTIFACT_HASH` field has been removed from the schema

File Structure

```
project/
|
├── db2_config.yaml      # Configuration file (REQUIRED)
├── db2_config.py        # Configuration loader
├── db2_evidence_base.py # Abstract base class
├── db2_evidence_ibm.py  # IBM DB implementation
├── db2_evidence_jdbc.py # JDBC implementation
├── db2_schema_ddl.sql   # Database schema DDL
└── README.md            # This file
```

Database Schema

The database uses the following structure:



Key Points:

- ARTIFACT table stores unique artifacts identified by ARTIFACT_PATH
- STEP_ARTIFACT junction table creates many-to-many relationships
- Same artifact can be used by multiple steps (no duplication)
- Constraint: UNIQUE (ARTIFACT_PATH) ensures artifact uniqueness

Database Setup

Before using the evidence loader, you must create the database and schema.

Step 1: Create Database (if needed)

On Linux/Unix where DB2 is installed:

```

# Switch to DB2 instance owner
su - db2inst1

# Create database
db2 create database DEPLOY using codeset UTF-8 territory en PAGESIZE 8192

# Verify database was created
db2 list database directory | grep DEPLOY
    
```

Step 2: Create Schema

Create the schema that will contain all tables:

```

-- Connect to database
db2 connect to DEPLOY user db2inst1

-- Verify connection
db2 "SELECT CURRENT SERVER FROM SYSIBM.SYSDUMMY1"

-- Create schema
db2 "CREATE SCHEMA DEPLOYZ AUTHORIZATION db2inst1"

-- Verify schema was created
db2 "SELECT SCHEMENAME FROM SYSCAT.SCHEMATA WHERE SCHEMENAME = 'DEPLOYZ'"
    
```

Step 3: Create Tables

Run the DDL script to create all tables:

```

# Execute the schema creation script
db2 -tvf db2_schema_ddl.sql
    
```

Step 4: Verify Tables

Check that all tables were created successfully:

```
-- List all tables in DEPLOYZ schema
db2 "SELECT TABNAME FROM SYSCAT.TABLES WHERE TABSCHEMA = 'DEPLOYZ' ORDER BY TABNAME"

-- Expected output:
-- TABNAME
-----
-- ACTION
-- ACTIVITY
-- ARTIFACT
-- DEPLOY
-- PROPERTIES
-- STEP
-- STEP_ARTIFACT
-- STEP_RESULT_DETAIL
-- V_ARTIFACT_USAGE
-- V_DEPLOYMENT_HIERARCHY
-- V_STEP_ARTIFACTS
```

Installation and Configuration

Option 1: IBM DB Driver (ibm_db)

Installation

```
pip install ibm_db
```

Configuration (Windows)

The IBM DB driver requires DB2 CLI libraries on Windows:

1. **Download DB2 CLI Driver**
 - Download from: [IBM Data Server Driver Package](#)
 - Extract to `C:\clidriver` (or another directory)

Useful Links

- **Official Documentation:** <https://github.com/ibmdb/python-ibmdb>
- **API Reference:** <https://github.com/ibmdb/python-ibmdb/wiki/APIs>
- **Installation Guide:** <https://github.com/ibmdb/python-ibmdb#installation>
- **Troubleshooting** (usefull for z/OS DB2 support): <https://github.com/ibmdb/python-ibmdb/blob/master/README.md#troubleshooting>

Prerequisites

- Python: 3.6+

Option 2: JDBC Driver (JayDeBeApi)

Installation

```
pip install JayDeBeApi
pip install JPyte1 # Required for JayDeBeApi
```

Configuration

1. **Download DB2 JDBC Driver**
 - Download `db2jcc4.jar` from: [IBM Data Server Driver for JDBC](#)
 - Or extract from DB2 installation: `<DB2_HOME>/java/db2jcc4.jar`
 - Save to an accessible directory (e.g., `/opt/drivers/db2jcc4.jar`)

Useful Links

- **JayDeBeApi Documentation:** <https://github.com/baztian/jaydebeapi>
- **JPyte1 Documentation:** <https://jpyte.readthedocs.io/>
- **DB2 JDBC Driver:** <https://www.ibm.com/support/pages/db2-jdbc-driver-versions-and-downloads>
- **DB2 JDBC Developer Guide:** <https://www.ibm.com/docs/en/db2/12.1.x?topic=apis-installing-data-server-driver-jdbc-sqlj>

- JayDeBeApi Examples: <https://github.com/baztian/jaydebeapi/blob/master/README.rst>

Prerequisites

- Java JDK: 8 or higher (JRE is sufficient)
- Python: 3.6+
- JAR File: db2jcc4.jar or db2jcc.jar

Usage

Configuration File

Update the `db2_config.yaml` file with your database connection settings:

```
# Driver selection: ibm_db or jdbc
driver: ibm_db

# Database schema
schema: DEPLOYZ

# IBM DB configuration (if driver: ibm_db)
ibm_db:
  database: DEPLOY
  hostname: localhost
  port: 50000
  protocol: TCPIP
  uid: ${DB2_USER}
  pwd: ${DB2_PASSWORD}

# JDBC configuration (if driver: jdbc)
jdbc:
  url: jdbc:db2://localhost:50000/DEPLOY
  username: ${DB2_USER}
  password: ${DB2_PASSWORD}
  driver_path: /opt/drivers/db2jcc4.jar
  driver_class: com.ibm.db2.jcc.DB2Driver

# Logging
logging:
  level: INFO
  format: '%(asctime)s - %(levelname)s - %(message)s'

# Options
options:
  verbose: true
  stop_on_error: false
```

Using Environment Variables in Configuration

You can use environment variables in your `db2_config.yaml`:

```
jdbc:
  url: jdbc:db2://${DB_HOST}:${DB_PORT}/${DB_NAME}
  username: ${DB_USER}
  password: ${DB_PASSWORD}
  driver_path: ${JDBC_DRIVER_PATH}
```

Then set the environment variables:

```

# Unix
export DB_HOST=localhost
export DB_PORT=50000
export DB_NAME=DEPLOY
export DB_USER=db2inst1
export DB_PASSWORD=secret
export JDBC_DRIVER_PATH=/opt/drivers/db2jcc4.jar

# Windows
set DB_HOST=localhost
set DB_PORT=50000
set DB_NAME=DEPLOY
set DB_USER=db2inst1
set DB_PASSWORD=secret
set JDBC_DRIVER_PATH=C:\drivers\db2jcc4.jar

```

Using IBM DB Driver

Command Line

```

# Set credentials
set DB2_USER=YOUR_DB2_USER
set DB2_PASSWORD=YOUR_DB2_PASSWORD

# Using default db2_config.yaml
python.exe db2_evidence_ibm.py evidences.yml

# Using custom config file
python.exe db2_evidence_ibm.py evidences.yml my_config.yaml

```

Using JDBC Driver

Command Line

```

# Set credentials
export DB2_USER=YOUR_DB2_USER
export DB2_PASSWORD=YOUR_DB2_PASSWORD

# Using default db2_config.yaml
python3 db2_evidence_jdbc.py evidences.yml

# Using custom config file
python3 db2_evidence_jdbc.py evidences.yml my_config.yaml

```

Querying Samples

```

-- List all deployments in a specific environment
SELECT DEPLOY_ID, ENVIRONMENT_NAME, DEPLOY_TIMESTAMP, STATUS
FROM DEPLOYZ.DEPLOY
WHERE ENVIRONMENT_NAME = 'PROD'
ORDER BY DEPLOY_TIMESTAMP DESC;

-- List all artifacts deployed in a specific environment
SELECT DISTINCT art.ARTIFACT_ID,
                art.ARTIFACT_NAME,
                art.ARTIFACT_TYPE,
                art.ARTIFACT_PATH
FROM DEPLOYZ.DEPLOY d
JOIN DEPLOYZ.ACTIVITY act ON d.DEPLOY_ID = act.DEPLOY_ID
JOIN DEPLOYZ.ACTION a ON act.ACTIVITY_ID = a.ACTIVITY_ID
JOIN DEPLOYZ.STEP s ON a.ACTION_ID = s.ACTION_ID
JOIN DEPLOYZ.STEP_STEPArtifact sa ON s.STEP_ID = sa.STEP_ID

```

```

JOIN DEPLOYZ.STEP_ARTIFACT sa ON s.STEP_ID = sa.STEP_ID
JOIN DEPLOYZ.ARTIFACT art ON sa.ARTIFACT_ID = art.ARTIFACT_ID
WHERE d.ENVIRONMENT_NAME = 'PROD'
ORDER BY art.ARTIFACT_NAME;

-- In which environments is the artifact named "LGACDB02" deployed?
SELECT DISTINCT
    d.ENVIRONMENT_NAME,
    d.DEPLOY_TIMESTAMP,
    act.ACTIVITY_NAME,
    a.ACTION_NAME,
    s.STEP_NAME,
    art.ARTIFACT_TYPE
FROM DEPLOYZ.DEPLOY d
JOIN DEPLOYZ.ACTIVITY act ON d.DEPLOY_ID = act.DEPLOY_ID
JOIN DEPLOYZ.ACTION a ON act.ACTIVITY_ID = a.ACTIVITY_ID
JOIN DEPLOYZ.STEP s ON a.ACTION_ID = s.ACTION_ID
JOIN DEPLOYZ.STEP_ARTIFACT sa ON s.STEP_ID = sa.STEP_ID
JOIN DEPLOYZ.ARTIFACT art ON sa.ARTIFACT_ID = art.ARTIFACT_ID
WHERE art.ARTIFACT_NAME = 'LGACDB02'
ORDER BY d.ENVIRONMENT_NAME, d.DEPLOY_TIMESTAMP

-- In which environments is the artifact named "LGACDB02" of type "CICSLOAD" deployed?
SELECT DISTINCT
    d.ENVIRONMENT_NAME,
    d.DEPLOY_TIMESTAMP,
    act.ACTIVITY_NAME,
    a.ACTION_NAME,
    s.STEP_NAME,
    art.ARTIFACT_TYPE
FROM DEPLOYZ.DEPLOY d
JOIN DEPLOYZ.ACTIVITY act ON d.DEPLOY_ID = act.DEPLOY_ID
JOIN DEPLOYZ.ACTION a ON act.ACTIVITY_ID = a.ACTIVITY_ID
JOIN DEPLOYZ.STEP s ON a.ACTION_ID = s.ACTION_ID
JOIN DEPLOYZ.STEP_ARTIFACT sa ON s.STEP_ID = sa.STEP_ID
JOIN DEPLOYZ.ARTIFACT art ON sa.ARTIFACT_ID = art.ARTIFACT_ID
WHERE art.ARTIFACT_NAME = 'LGACDB02' and art.ARTIFACT_TYPE = 'CICSLOAD'
ORDER BY d.ENVIRONMENT_NAME, d.DEPLOY_TIMESTAMP

-- List the properties (PROPERTIES) of an activity for a deployment with a specific DEPLOY_TIMESTAMP
SELECT
    act.ACTIVITY_ID,
    act.ACTIVITY_NAME,
    p.PROPERTY_KEY,
    p.PROPERTY_VALUE
FROM DEPLOYZ.DEPLOY d
JOIN DEPLOYZ.ACTIVITY act ON d.DEPLOY_ID = act.DEPLOY_ID
JOIN DEPLOYZ.PROPERTIES p
    ON p.ENTITY_TYPE = 'ACTIVITY'
    AND p.ENTITY_ID = act.ACTIVITY_ID
WHERE d.DEPLOY_TIMESTAMP = TIMESTAMP('2025-11-27 12:26:21.0')
ORDER BY act.ACTIVITY_ID, p.PROPERTY_KEY;

-- List the properties of a specific artifact (by name and type) for a given deployment timestamp
SELECT
    d.DEPLOY_ID,
    d.ENVIRONMENT_NAME,
    d.DEPLOY_TIMESTAMP,
    art.ARTIFACT_ID,
    art.ARTIFACT_NAME,
    art.ARTIFACT_TYPE,
    p.PROPERTY_KEY,
    p.PROPERTY_VALUE
FROM DEPLOYZ.DEPLOY d
JOIN DEPLOYZ.ARTIFACT art
    ON EXISTS (

```

```

SELECT 1
FROM DEPLOYZ.STEP_ARTIFACT sa
JOIN DEPLOYZ.STEP s ON sa.STEP_ID = s.STEP_ID
JOIN DEPLOYZ.ACTION a ON s.ACTION_ID = a.ACTION_ID
JOIN DEPLOYZ.ACTIVITY act ON a.ACTIVITY_ID = act.ACTIVITY_ID
WHERE sa.ARTIFACT_ID = art.ARTIFACT_ID
AND act.DEPLOY_ID = d.DEPLOY_ID
)
JOIN DEPLOYZ.PROPERTIES p
ON p.ENTITY_TYPE = 'ARTIFACT'
AND p.ENTITY_ID = art.ARTIFACT_ID
AND p.DEPLOY_ID = d.DEPLOY_ID
WHERE d.DEPLOY_TIMESTAMP = TIMESTAMP('2025-11-27 12:26:21.0') and art.ARTIFACT_NAME = 'LGACDB02'
ORDER BY art.ARTIFACT_ID, p.PROPERTY_KEY;

```

Driver Comparison

| Feature | IBM DB (<code>ibm_db</code>) | JDBC (<code>JayDeBeApi</code>) |
|---------------|--------------------------------|----------------------------------|
| Performance | ★★★★ Excellent | ★★★★ Very Good |
| Installation | ★★ Medium (requires CLI) | ★★★★ Easy (just JAR) |
| Portability | ★★★ OS-dependent | ★★★★ Cross-platform |
| Compatibility | ★★★★ Native DB2 | ★★★★ Standard JDBC |
| Windows | ★★ Complex setup | ★★★★ Simple |
| Linux | ★★★★ Native | ★★★★ Requires Java |
| z/OS | ★★★★ Native For DB2 on z/OS | ★★★★ Requires Java |

When to Use Which Driver?

Use `ibm_db` if:

- DB2 CLI is already installed
- Maximum performance required
- Running on Linux/AIX
- Running on z/OS to access DB2 on z/OS
- Need advanced DB2 features

Use `JayDeBeApi` if:

- Simplified installation desired
- Cross-database portability needed
- Docker/containerized deployment

Additional Resources

DB2 Documentation

- [DB2 Knowledge Center](#)

License